
AI Planning

1. Bellman Equation

- Apply concept of dynamic programming and optimal substructure to do **policy evaluation**

$$V(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) V(s')$$

- s: Any given state
- s': Any possible next state
- a: Any possible action at state s
- V(s): Value of state
- R(s); Reward of state
- A(s): Set of all possible actions at state s
- γ : Exploration Factor
- P(s'|s,a): Probability of s' given (s,a)
- V(s'): Value of s'
- **Value of a state = the state's value + Exploration Factor x the value of the best action at that state**
 - What is the value of the best action?
 - Given (s,a), Sum all possible next states x the probability of reaching those states.

Bellman Equation Variants

R(s,a,s')

$$V(s) = \max_{a \in A} \sum_{s' \in S} T(s, a, s') [R(s, a, s') + \gamma V(s')]$$

- Reward for transitioning from s to s' via a

- Choose the action that has gives the highest reward (based on transition), and discounted future reward for reaching that state

R(s,a)

$$V(s) = \max_{a \in A} R(s, a) + \gamma \sum_{s' \in S} T(s, a, s') V(s')$$

- Reward for taking action a at state s
- Choose the action that has gives the highest reward, and also factor in the discounted next state's reward

Converting between reward functions

$$R'(s, a) = \sum_{s' \in S} T(s, a, s') R(s, a, s')$$

- The reward for taking an action is the sum of (all its possible transitions) x (related transition rewards)

$$R'(post(s, a)) = \gamma^{-1/2} R(s, a)$$

- s' is abstracted out to "post(s,a)", which refers to the "post-state" for every (s,a). In this R:

```
// (s,a) always goes to post(s,a)
T'(s,a,post(s,a)) = 1
// The probability is factored in here
T'(post(s,a),b,s') = T(s,a,s')
// The states themselves don't have rewards
R'(s) = 0
// The reward for taking the action is put into the pseudo "post-state"
R'(post(s,a)) = gamma^(-1/2)R(s,a)
gamma' = gamma^(1/2)
```

2. Value Iteration Algorithm

- Repeatedly update the Utility function U(s) with the bellman equation / bellman update above.

$$U_{t+1}(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U_t(s')$$

- Notice that U is only updated after every single state has been looped through.

```
// Returns Utility function U(s)
function value-Iteration(mdp, err_threshold)
  loop:
    U_t = U_t+1, max_delta = 0
    for all states s in S:
      Update U_t+1 with bellman equation.
      Update max_delta if abs(U(s) - U_t+1(s)) exceed max_delta
      break loop if max_delta less than err_threshold(1-gamma)/gamma
  return U_t
```

3. Policy Evaluation Algorithm

- Same as Value Iteration.
 - But we want to **evaluate a given policy π_i** . Thus, **we don't need to take the max, we just need to use the policy.**
 - Hence: Instead of taking the best action, take the action that the policy thinks is the best action

$$U_i(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi_i(s)) U_i(s')$$

4. Policy Iteration Algorithm

- We don't need to calculate U(s) so accurately if we just want to find the optimal policy
- Two step approach:
 1. **Init:** Start with initial policy π_0
 2. **Policy Evaluation:** Calculate Utility function U_i given current policy. See Policy Evaluation above.
 3. **Policy Improvement:** Calculate new policy π_{i+1} using U_i .

4. **Termination:** Repeat until the new and old policy are the same (no change).

```
// Returns policy pi
function Policy-Iteration(mdp)
    U = set all to 0
    pi = random policy
    loop
        unchanged = true
        U_i = Policy-Evaluation(pi, U, mdp)
        for all states s in S:
            If the best action at s is different:
                Update pi[s]
                unchanged = false
        break if unchanged
    return pi
```

The check for the best action is done as follows:

$$\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U_i(s') > \sum_{s'} P(s' | s, \pi_i(s)) U_i(s')$$

- Identify the action with highest utility at s
- Check if it is the same action taken in the policy

5. Modified Policy Iteration

- Only do k iterations instead of until no change

6. Model-Based Prediction

- Prediction: Given traces of a policy and the final reward, learn the utility function by constructing the model from the data using an agent:
 - Learn **transition model** and **reward function** using an **Adaptive Dynamic Programming (ADP)** agent.
 - **Calculate the Utility function.**

```
// persistent variables
```

```

s_prev

pi: policy

mdp: current constructed model, rewards and discount

U: table of utilities

Nsa: Table keeping count of number of times at state s, action a was
    taken

Nsas': Table keeping count of number of times the next state was s'
    given (s,a) (thus s,a,s')

// Called everytime a new percept is observed by the agent
// Percept: (previous state s, current state s', current reward r')
function PASSIVE-ADP-AGENT(percept)
    if s_curr is new:
        U[s_curr] = r_curr
        R[s_curr] = r_curr
    if s_prev not null:
        // Increment counters
        Nsa[s_prev,a]++
        Ns'sa[s_curr,s_prev,a]++

        // Update transition function
        // for every known reachable state by (s,a)
        for every state where Ns'sa[state,s,a] > 0:
            // # of times state happened when action was taken
            T(s,a,s') = Ns'as[s,a,s'] / Nsa[s,a]
    U = Policy-Evaluation(pi, U, mdp)
    if s_curr is terminal:
        s_prev,a = null
    else
        s_prev,a = s_curr, pi[s_curr]
    return a

```

7. Model-Based Control

- Learn the policy, not the utility
- Just replace Policy-Evaluation with Policy-Iteration
- The policy no longer stays fixed but changes as transitions and rewards learnt
- Note that however the algorithm is greedy and may not return the optimal value

- ## 8. Monte Carlo Learning

- Assume state s encountered k times with k returns, and each summed reward is stored in $G_i(s)$.

- $$U_k(s) = \frac{1}{k} \sum_{i=1}^k G_i(s)$$

6

$$U_k(s) = \frac{1}{k} G_k(s) + \frac{1}{k} * (k-1) U_{k-1}(s)$$

$$U_k(s) = \frac{k}{k} U_{k-1}(s) - \frac{1}{k} U_{k-1}(s) + \frac{1}{k} G_k(s)$$

$$U_k(s) = U_{k-1}(s) + \frac{1}{k} (G_k(s) - U_{k-1}(s))$$

- The difference between the current $U_k(s)$ and the previous $U(s)$ is the prediction error (if you want to minimize absolute loss, you can use median instead of average)
- Note that $U(s)$ here can also apply to $U(s,a)$, and can also be renamed as $Q(s)$ and $Q(s,a)$

```
# Monte Carlo Prediction (Learning U(s))
After every trial:
    Take every state (or first state occurrence) and calculate its value
    sum
    Increment Ns
    U_k(s) = U_k-1(s) + 1/k(sum - U_k-1(s))

# Monte Carlo Control (Learning pi)
After every trial:
    Take every (s,a):
        newR(s,a) = take every/1st (s,a) and calculate value sum
        Increment Ns
        Q(s,a) = Q_k-1(s,a) + 1/k(newR(s,a) - Q_k-1(s,a))
    For every s:
        pi(s) = Take action that maximizes Q(s,a)
```

Advantages

- Simple
- Unbiased estimate

Disadvantages

- Must wait until full trial is done in order to perform learning

- High variance (reward is the sum of many rewards along the trial), so need many trials to get it right

9. Temporal Difference Learning

- Very similar to Monte Carlo Learning
- Replace the

$$U^\pi(s) = U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$

- Transition from state s to s' .
- Alpha is the learning rate.
 - Converges if alpha decreases with the number of times the state has been visited (think GLIE)
- $R(s) + \gamma U^\pi(s')$: **Temporal Difference target**
 - (this is basically bellman update, policy evaluation)
- $R(s) + \gamma U^\pi(s') - U^\pi(s)$: **Temporal Difference error**
 - (amount of utility change from previous state s to current state s')
 - difference between the estimated reward at any given state or time step and the actual reward received
 - Utility of current state + discounted future reward(future state) - expectedRewardWithIncludesFutureRewards(s)

```
# Prediction
TD-Agent:
for every percept (curr_state s', immediate reward r')
if s' is new:
    U[s'] = r'
if s not null:
    NS[s]++
    U[s] = U[s] + alpha*NS[s]*(r+gamma*U[s']-U[s])
if s' terminal:
    s,a,r = null
else: # update prev data
    s,a,r = s',pi[s'],r'
```



```
return a
```

Advantages

- Can learn with every step (in a trial)
- Usually converges faster in practice

Disadvantages

- Online; lower variance, but estimate on how good your estimate is; biased
- Assumes MDP
- Error can go in any direction

SARSA uses TD learning w.r.t. a specific policy:

$$U^\pi(s, a) = U^\pi(s, a) + \alpha(R(s, a) + \gamma U^\pi(s', \pi(s')) - U^\pi(s, a))$$

Q-Learning uses TD learning w.r.t. the optimal policy (s' is the next state after (s,a) is taken):

$$Q(s, a) = Q(s, a) + \alpha(R(s, a) + \gamma \max_{a \in A(s')} Q(s', a) - Q(s, a))$$

10. n-step TD, TD(lambda)

Alister Reis has a very good blog post on this <https://amreis.github.io/ml/reinforce/2017/11/02/reinforcement-learning-eligibility-traces.html>

